

# A Cloud-based Application Storage Service for Ibis

Bas Boterman

June 25, 2009



# Outline

Introduction

Project Details

Benchmarks

Applications

Conclusion

# Background

- ▶ April 7, 2008, the Google App Engine, a platform for developing and hosting web applications, was introduced

# Background

- ▶ April 7, 2008, the Google App Engine, a platform for developing and hosting web applications, was introduced
- ▶ To what extent can we use the Google App Engine for scientific purposes?

# Background

- ▶ April 7, 2008, the Google App Engine, a platform for developing and hosting web applications, was introduced
- ▶ To what extent can we use the Google App Engine for scientific purposes?
- ▶ Robust distributed datastore with transactions and queries

# Background

- ▶ April 7, 2008, the Google App Engine, a platform for developing and hosting web applications, was introduced
- ▶ To what extent can we use the Google App Engine for scientific purposes?
- ▶ Robust distributed datastore with transactions and queries
- ▶ This project is about designing an application storage server for Ibis

# Thesis Title

## A Cloud-based Application Storage Service for Ibis

# Thesis Title

## A **Cloud-based** Application Storage Service for Ibis



# Thesis Title

## A **Cloud-based** Application Storage Service for Ibis

*“A style of computing in which dynamically scalable resources are provided as a service over the Internet.”*

E.g. the Google App Engine.

# Thesis Title


A Cloud-based **Application Storage Service** for Ibis

# Thesis Title

## A Cloud-based **Application Storage Service** for Ibis

*"A service used by applications to store and retrieve application data."*

Example:



# Thesis Title

A Cloud-based Application Storage Service for **Ibis**

# Thesis Title

## A Cloud-based Application Storage Service for **Ibis**

*“The main goal of the Ibis project is to create an efficient Java-based platform for grid computing.”*

- ▶ JavaGAT (Grid Application Toolkit)
- ▶ IPL (Ibis Portability Layer)

# Google App Engine

*“A platform for building and hosting web applications on the Google infrastructure.”*

- ▶ Python programming language
- ▶ Web application framework
- ▶ Authentication through Google Accounts
- ▶ 10 GB traffic per day
- ▶ 1 GB of data storage
- ▶ Free of charge.

# (Dis)Advantages of the Google App Engine

## Advantages:

- ▶ Powerful (distributed) database with query engine and transactions
- ▶ Replicating servers to improve scalability
- ▶ Authentication through Google Accounts
- ▶ Built-in frameworks (Django, WebOb, PyYAML)

# (Dis)Advantages of the Google App Engine

## Advantages:


- ▶ Powerful (distributed) database with query engine and transactions
- ▶ Replicating servers to improve scalability
- ▶ Authentication through Google Accounts
- ▶ Built-in frameworks (Django, WebOb, PyYAML)

## Disadvantages:

- ▶ Python 2.5 (limited functionality, compared to 3.0)
- ▶ HTTP(S) only
- ▶ Sandbox (no `fork()`, no threads, no file system)



# Project Overview



# Server Design


Public functions:

- ▶ `add(data, metadata, path)`
- ▶ `del(path)`
- ▶ `get(path)`
- ▶ `getmd(path)`
- ▶ `find(metadata)`
  
- ▶ `data`: Text
- ▶ `metadata`: Key-value pairs of Strings
- ▶ `path`: String


# Client Library

- ▶ Written in Java
- ▶ Implements the Google App Engine Advert server's protocol
  - ▶ Base64 encoding
  - ▶ JSON encoding (XML alternative)
  - ▶ Error handling
- ▶ Establishes connections over HTTP(S)
- ▶ Takes care of authentication
  - ▶ Authentication through *Google ClientLogin*
  - ▶ Automatically refresh authentication cookie


# Adding Objects of Variable Size




## Adding Objects with Variable Meta Data




## Retrieving Objects with Variable Meta Data



# Finding Meta Data



# Performing Operations in Parallel





# Benchmark Evaluation

- ▶ Completion time depends on object size/meta data size, and on number of clients connecting simultaneously
- ▶ Completion time is independent on the datastore size
- ▶ Google has some restrictions
  - ▶ Response time cannot be over 30 seconds
  - ▶ Quota temporarily exceeds with requests larger than 5 MB
  - ▶ No more than 500 data items can be manipulated in one call
  - ▶ No more than 100 clients can use ClientLogin simultaneously

# Outline

Introduction

Project Details

Benchmarks

**Applications**


Conclusion

# JavaGAT

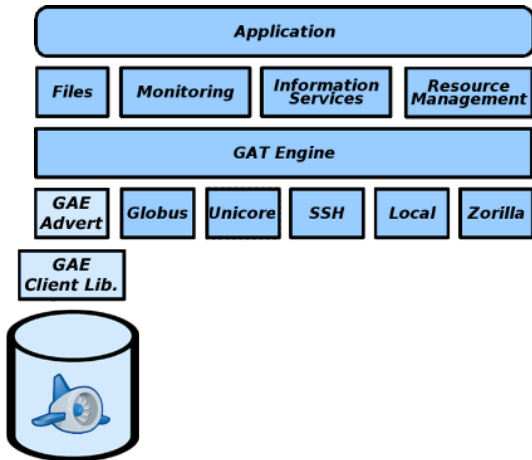
*JavaGAT offers a set of coordinated, generic and flexible APIs for accessing grid services from application codes, portals, data managements systems, etc.*

- ▶ File operations
- ▶ Job submission
- ▶ Monitoring
- ▶ Access to information services (AdvertService)

# JavaGAT structure



# JavaGAT structure with App Engine AdvertService Adaptor



# App Engine AdvertService Adaptor

- ▶ So far, only a local AdvertService existed
- ▶ Scalable design (local AdvertService does not scale)
- ▶ Completely transparent to user
- ▶ Path handling is done locally

# IPL

*“The IPL is a communication library which is specifically designed for usage in a grid environment.”*

- ▶ Communication library for distributed applications
- ▶ All applications discover each other through an IPL registry server
- ▶ IPL registry server address needs to be passed to each Ibis application

## IPL Example

```
$ $IPL_HOME/bin/ipl-server --events
```

Ibis server running on 130.37.20.18-8888~bbn230

List of Services:

Bootstrap service on virtual port 303

Central Registry service on virtual port 302

Management service

Known hubs now: 130.37.20.18-8888~bbn230

```
$ $IPL_HOME/bin/ipl-run \  
-Dibis.server.address=130.37.20.18-8888~bbn230 \  
-Dibis.pool.name=test \  
ibis.ipl.examples.Hello
```



## IPL Example Using Advert Server

```
$ $IPL_HOME/bin/ipl-server --events \  
--advert google://jondoe.appspot.com/identifier \  
--user jondoe@gmail.com --pass north23AZ \  
--metadata author=jondoe,created=24jun,pool=test,color=purple
```

Ibis server running on 130.37.20.18-8888~bbn230

List of Services:

- Bootstrap service on virtual port 303

- Central Registry service on virtual port 302

- Management service

Known hubs now: 130.37.20.18-8888~bbn230

## IPL Example Using Advert Server

```
$ $IPL_HOME/bin/ipl-run \  
-Dibis.advert.address=google://jondoe.appspot.com \  
-Dibis.advert.username=jondoe \  
-Dibis.advert.password=north23AZ \  
-Dibis.advert.metadata=created=24jun,color=purple \  
-Dibis.pool.name=test \  
ibis.ipl.examples.Hello
```

# IPL Server Bootstrap Mechanism

- ▶ No mechanism for bootstrapping IPL servers existed
- ▶ Large range of IPL servers can be started, whilst only remembering one (advert) address
- ▶ Different groups of Ibises can be started without (manually) having to configure each Ibis

# Conclusion

## Contributions:

- ▶ A Google App Engine Advert server, written in Python
- ▶ An Advert client library, written in Java
- ▶ An App Engine AdvertService adaptor for JavaGAT
- ▶ An IPL Server bootstrap mechanism

# Conclusion

## Contributions:

- ▶ A Google App Engine Advert server, written in Python
- ▶ An Advert client library, written in Java
- ▶ An App Engine AdvertService adaptor for JavaGAT
- ▶ An IPL Server bootstrap mechanism

## Properties:

- ▶ Scalable application and datastore
- ▶ High bandwidth (until quota reached)
- ▶ No guarantees

# Conclusion

## Contributions:

- ▶ A Google App Engine Advert server, written in Python
- ▶ An Advert client library, written in Java
- ▶ An App Engine AdvertService adaptor for JavaGAT
- ▶ An IPL Server bootstrap mechanism

## Properties:

- ▶ Scalable application and datastore
- ▶ High bandwidth (until quota reached)
- ▶ No guarantees

## Future work:

- ▶ An App Engine Advert server, written in Java